


Deep Learning & Transfer Learning

A Comparative Analysis of Image Classification on CIFAR-10


Baseline



Deep Scaling



Transfer Learning



Airplane



Automobile



Bird



Cat



Deer



Dog



Frog



Horse



Ship



Truck


The Baseline: Simple CNN

Model Stats

Total Parameters: 315,722



Accuracy: ~64% - 68%

Structure: 2 Conv Blocks + Dense





Layer (type)	Output Shape	Param #
conv2d_140 (Conv2D)	(None, 30, 30, 32)	896
max_pooling2d_84 (MaxPooling2D)	(None, 15, 15, 32)	0
conv2d_141 (Conv2D)	(None, 13, 13, 64)	18,496
conv2d_141 (Conv2D)	(None, 13, 13, 64)	18,496
max_pooling2d_85 (MaxPooling2D)	(None, 6, 6, 64)	0
flatten_38 (Flatten)	(None, 2304)	0
dense_76 (Dense)	(None, 128)	295,840
dense_77 (Dense)	(None, 10)	1,290


Total params: 315,722 (1.28 MB)
Trainable params: 315,722 (1.29 MB)
Non-trainable params: 0 (0.00 B)





Escalation: Deep CNN Architecture




Deeper CNN Architecture




deep_baseline(650K): F1=32, F2=64, F3= 128, F4= 256, Dense=256, Batch=32, Epochs= 50 46 LR= 0.001± 0.00003



deep_small(165K) : F1=32, F2=64, F3= 128, F4= 256, Dense=256, Batch=32, Epochs=50-35, LR = 0.0005 0.00001




Simple CNN vs Deeper CNN



Transfer Learning: Pretrained Models

EfficientNet B0



EfficientNetB0

Retrain almost 0% of the layers:

Layer (type)	Output Shape	Param #
input_layer_20 (InputLayer)	(None, 32, 32, 3)	0
resizing_10 (Resizing)	(None, 240, 240, 3)	0
lambda_10 (Lambda)	(None, 240, 240, 3)	0
efficientnetb0 (Functional)	(None, 7, 7, 1280)	4,049,571
global_average_pooling2d_10 (GlobalAveragePooling2D)	(None, 1280)	0
dense_20 (Dense)	(None, 128)	163,968
dense_21 (Dense)	(None, 10)	1,290

Total params: 4,214,829 (16.08 MB)

Trainable params: 165,258 (645.54 KB)

Non-trainable params: 4,049,571 (15.45 MB)

accuracy:
~92 %

Training time:
7min 46s

- Same architecture
- difference in “Trainable params”, training time and accuracy

retrain 20% of the layers:

Layer (type)	Output Shape	Param #
input_layer_68 (InputLayer)	(None, 32, 32, 3)	0
resizing_30 (Resizing)	(None, 240, 240, 3)	0
lambda_29 (Lambda)	(None, 240, 240, 3)	0
efficientnetb0 (Functional)	(None, 7, 7, 1280)	4,049,571
global_average_pooling2d_33 (GlobalAveragePooling2D)	(None, 1280)	0
dense_66 (Dense)	(None, 128)	163,968
dense_67 (Dense)	(None, 10)	1,290

Total params: 4,214,829 (16.08 MB)

Trainable params: 2,470,554 (9.42 MB)

Non-trainable params: 1,744,275 (6.65 MB)


Training time:
11min 43s

accuracy:
~95 %

EfficientNetB0


retrain 20% of the layers:

	precision	recall	f1-score	support
airplane	0.96	0.95	0.95	1000
automobile	0.97	0.97	0.97	1000
bird	0.93	0.94	0.94	1000
cat	0.88	0.91	0.89	1000
deer	0.95	0.94	0.95	1000
dog	0.92	0.89	0.91	1000
frog	0.96	0.98	0.97	1000
horse	0.96	0.96	0.96	1000
ship	0.97	0.97	0.97	1000
truck	0.96	0.96	0.96	1000
accuracy			0.95	10000
macro avg	0.95	0.95	0.95	10000
weighted avg	0.95	0.95	0.95	10000




EfficientNetB0

retrain 20% of the layers:




Predicting the class of our own pet

Original picture 3605x3605 px:




Resized picture 32x32 px:



1/1 ————— 0s 48ms/step
predicted class: cat

Conclusion: Performance vs. Cost



Final Verdict:

Transfer Learning achieves a decisive victory in accuracy (95%), offering the best performance-to-effort ratio for general classification.

However, this comes with a 7M parameter footprint and increased pre-processing overhead.